

OH-MY-
CLAUDECODE:
MULTI-AGENT
ORCHESTRATION
FOR CLAUDE CODE

ZERO LEARNING CURVE. MAXIMUM POWER.

[Speaker Name]

Version 3.6.3

AGENDA

Time	Topic
------	-------

0:00	What is OMC?
------	--------------

0:10	The 5 Key Execution Modes
------	---------------------------

0:30	The Agent System
------	------------------

0:40	Live Demo Scenarios
------	---------------------

0:48	Developer Experience
------	----------------------

0:54	Getting Started
------	-----------------

0:58	Q&A
------	-----

THE PROBLEM

Developers today face:

THE PROBLEM

Developers today face:

- Manual coordination of complex multi-step tasks

THE PROBLEM

Developers today face:

- Manual coordination of complex multi-step tasks
- Constant context-switching between different concerns

THE PROBLEM

Developers today face:

- Manual coordination of complex multi-step tasks
- Constant context-switching between different concerns
- Single-threaded AI interactions that don't scale

THE PROBLEM

Developers today face:

- Manual coordination of complex multi-step tasks
- Constant context-switching between different concerns
- Single-threaded AI interactions that don't scale
- No persistence - AI gives up when tasks get hard

THE PROBLEM

Developers today face:

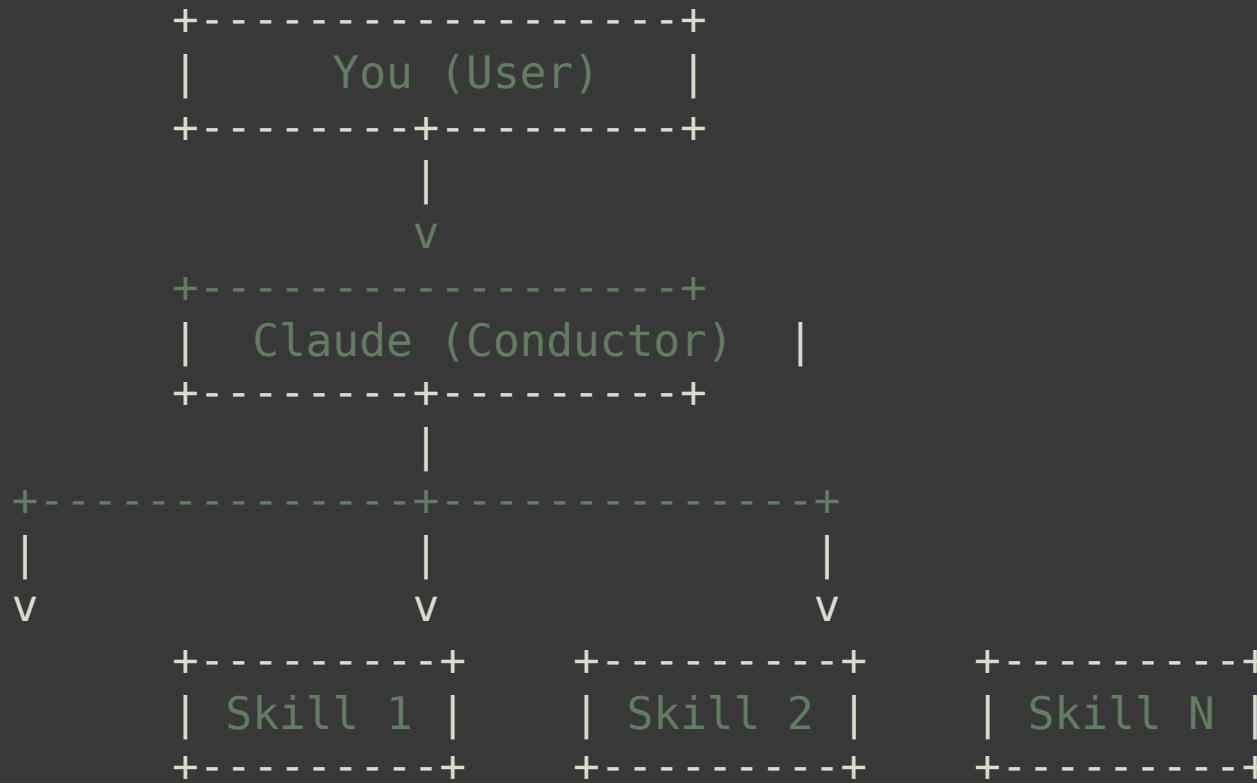
- Manual coordination of complex multi-step tasks
- Constant context-switching between different concerns
- Single-threaded AI interactions that don't scale
- No persistence - AI gives up when tasks get hard
- Token waste - using expensive models for simple tasks

SECTION 1

WHAT IS OMC?

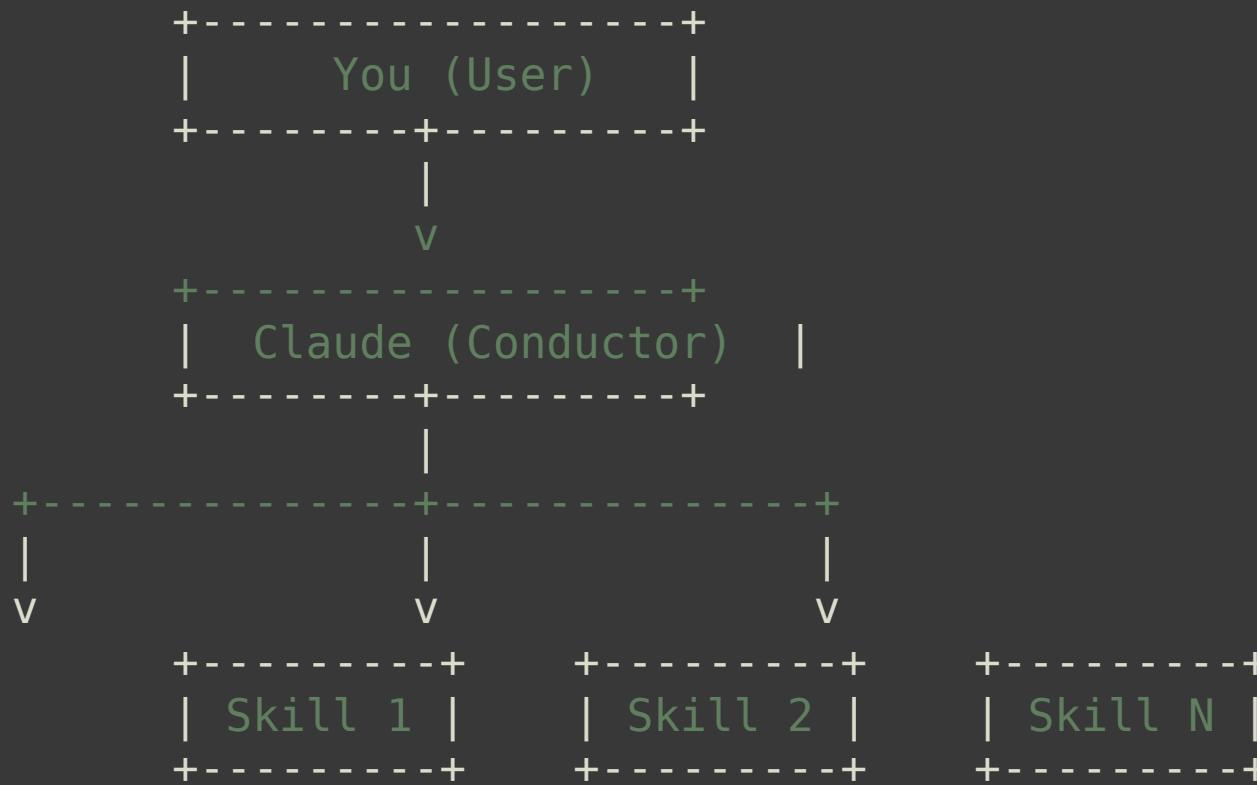
WHAT IS OH-MY-CLAUDECODE?

A multi-agent orchestration system for Claude Code



WHAT IS OH-MY-CLAUDECODE?

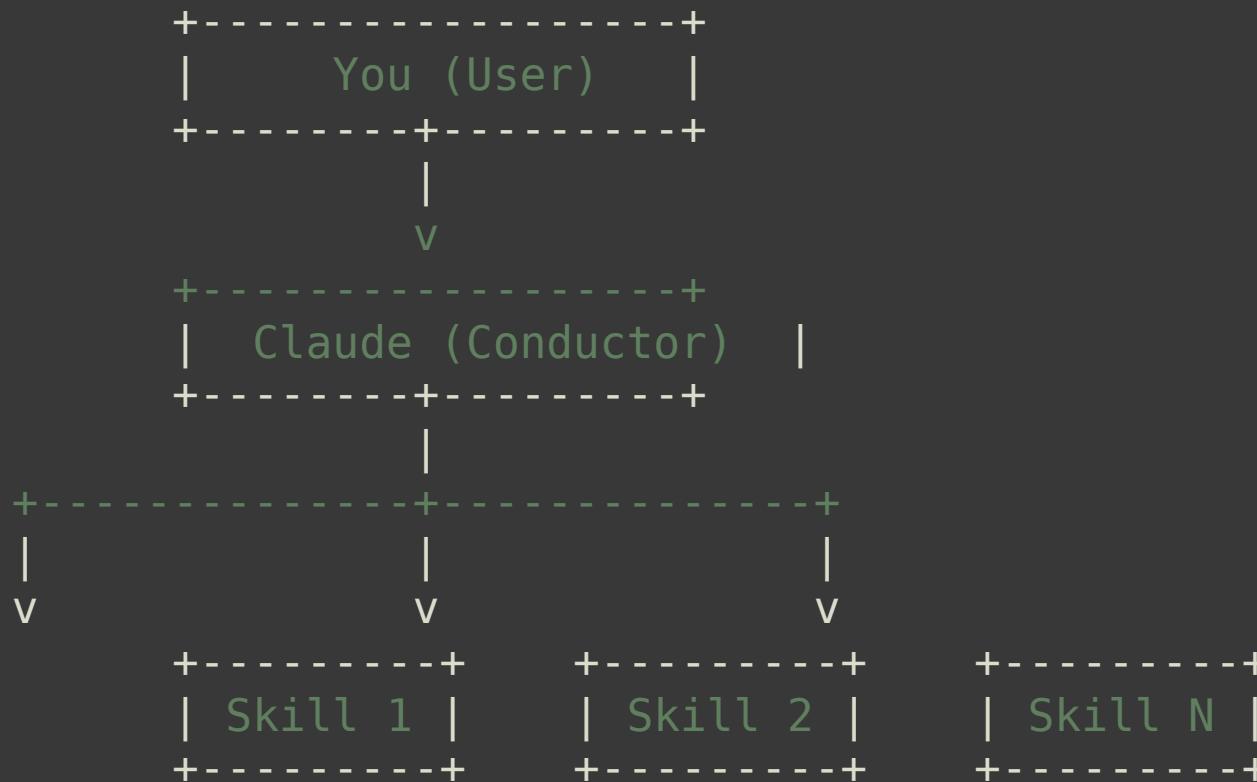
A multi-agent orchestration system for Claude Code



- 32 specialized agents

WHAT IS OH-MY-CLAUDECODE?

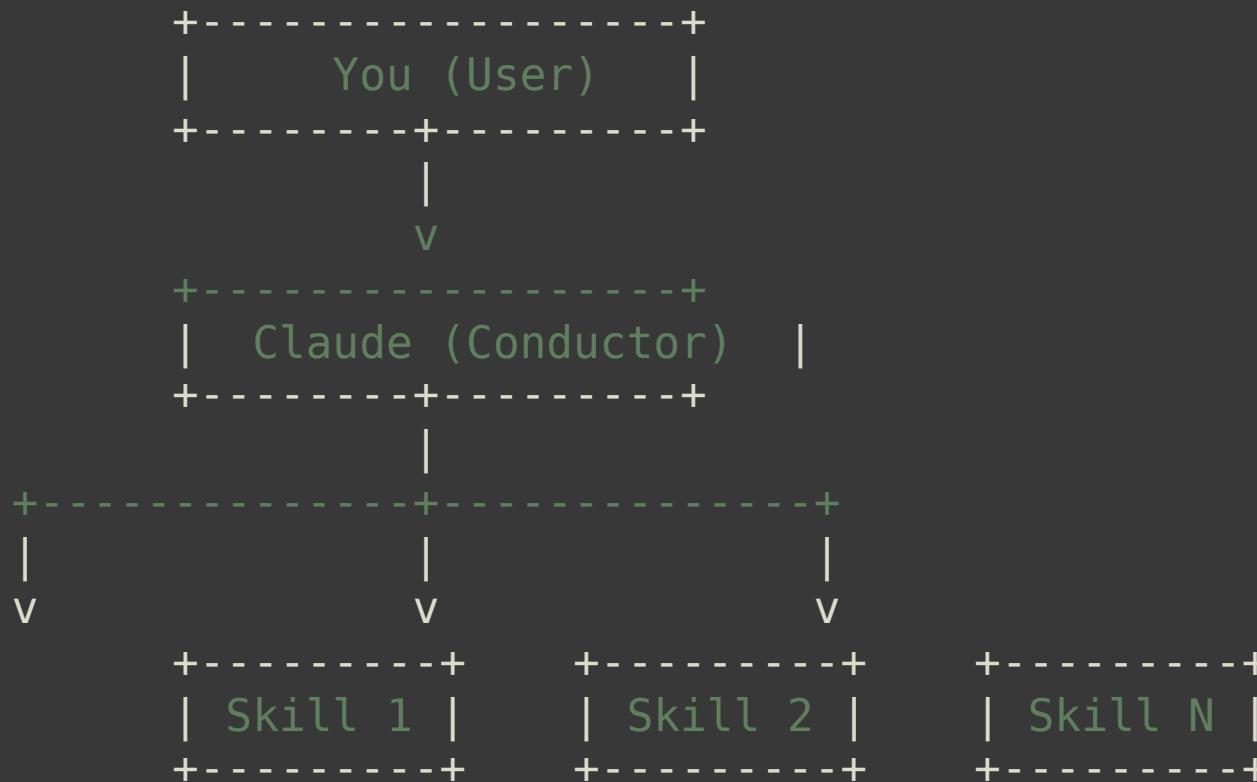
A multi-agent orchestration system for Claude Code



- 32 specialized agents
- 35+ skills

WHAT IS OH-MY-CLAUDECODE?

A multi-agent orchestration system for Claude Code



- 32 specialized agents
- 35+ skills
- Zero configuration required

THE PHILOSOPHY

"You are a CONDUCTOR, not a performer."

Traditional AI Workflow:

User -> Claude -> [Does everything itself]

OMC Workflow:

User -> Claude (Conductor) -> [Delegates to specialists]



architect
(analysis)

executor
(implementation)

designer
(UI/UX)

Claude becomes an intelligent orchestrator that delegates to the right specialist for each task.

BEFORE VS AFTER OMC

Aspect	Before OMC	After OMC
Task execution	Single-threaded	Parallel agents
Complex tasks	Manual breakdown	Automatic decomposition
Model selection	Always same model	Smart routing (Haiku/Sonnet/Opus)
Persistence	Gives up easily	Continues until verified

Aspect	Before OMC	After OMC
Cost	Expensive	30-50% savings
Learning curve	Command memorization	Natural language

Example - "Fix all TypeScript errors":

Before: You manually find and fix each error sequentially

After: 5 parallel agents claim and fix errors simultaneously

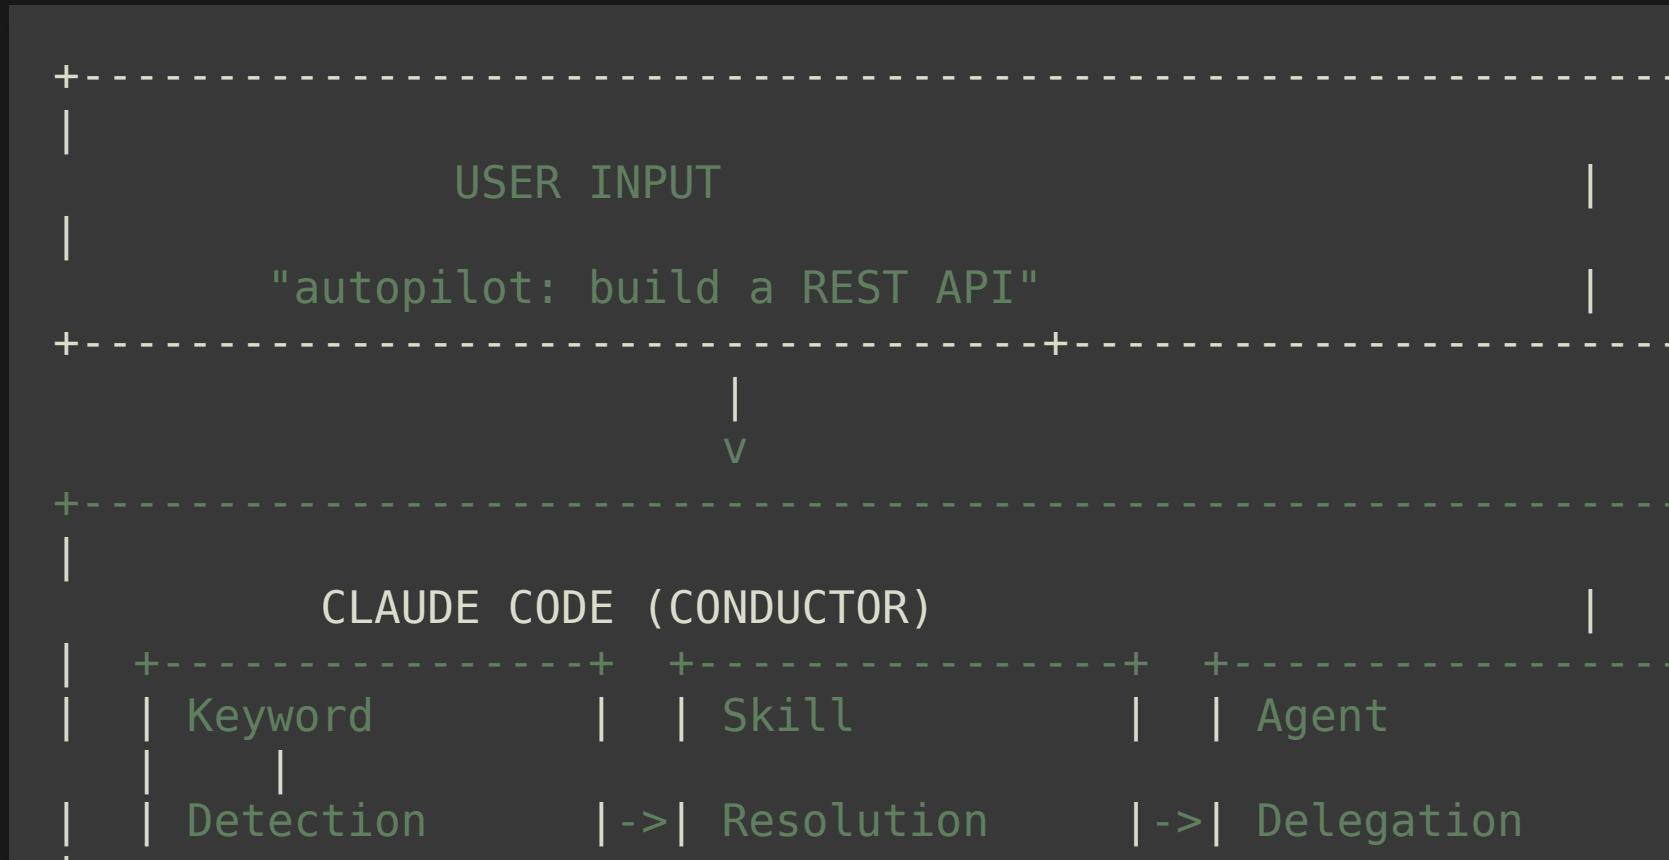
KEY STATISTICS

Metric	Value
Specialized Agents	32
Skills	35+
Execution Modes	8
Lifecycle Hooks	19
Model Tiers	3 (Haiku, Sonnet, Opus)
License	MIT

Token Cost Comparison:

Model	Input	Output
Haiku	\$0.25/1M	\$1.25/1M
Sonnet	\$3/1M	\$15/1M
Opus	\$15/1M	\$75/1M

ARCHITECTURE OVERVIEW



SECTION 2

THE 5 KEY EXECUTION MODES

MODE 1: AUTOPILOT - WHAT IS IT?

Full autonomous execution from idea to working code

```
"autopilot: build a REST API for a bookstore"
```

5 Phases:

1. **Expansion** - Turn vague idea into detailed spec
2. **Planning** - Create implementation plan with validation
3. **Execution** - Build with parallel agents (Ralph + Ultrawork)
4. **QA** - Test until everything passes (up to 5 cycles)
5. **Validation** - Multi-reviewer approval (Architect + Security + Code Review)

MODE 1: AUTOPILOT - HOW IT WORKS

Phase 0: EXPANSION

|
+--> Analyst (Opus) extracts requirements
+--> Architect (Opus) creates technical spec

|
v

Phase 1: PLANNING

|
+--> Architect creates plan (direct mode)
+--> Critic validates plan

|
v

Phase 2: EXECUTION

|
+--> Ralph + Ultrawork activated

MODE 1: AUTOPILOT - WHEN TO USE IT

Best For:

- New projects from scratch
- Complete feature implementations
- End-to-end workflows

Trigger Keywords:

```
autopilot, auto pilot, autonomous  
build me, create me, make me  
full auto, handle it all  
I want a/an...
```

Example Commands:

```
autopilot: build a REST API with CRUD for inventory  
/oh-my-claudecode:autopilot Add OAuth2 authentication  
autopilot: create a CLI tool that tracks daily habits
```

MODE 2: ULTRAPILOT - WHAT IS IT?

Parallel autopilot with up to 5 concurrent workers

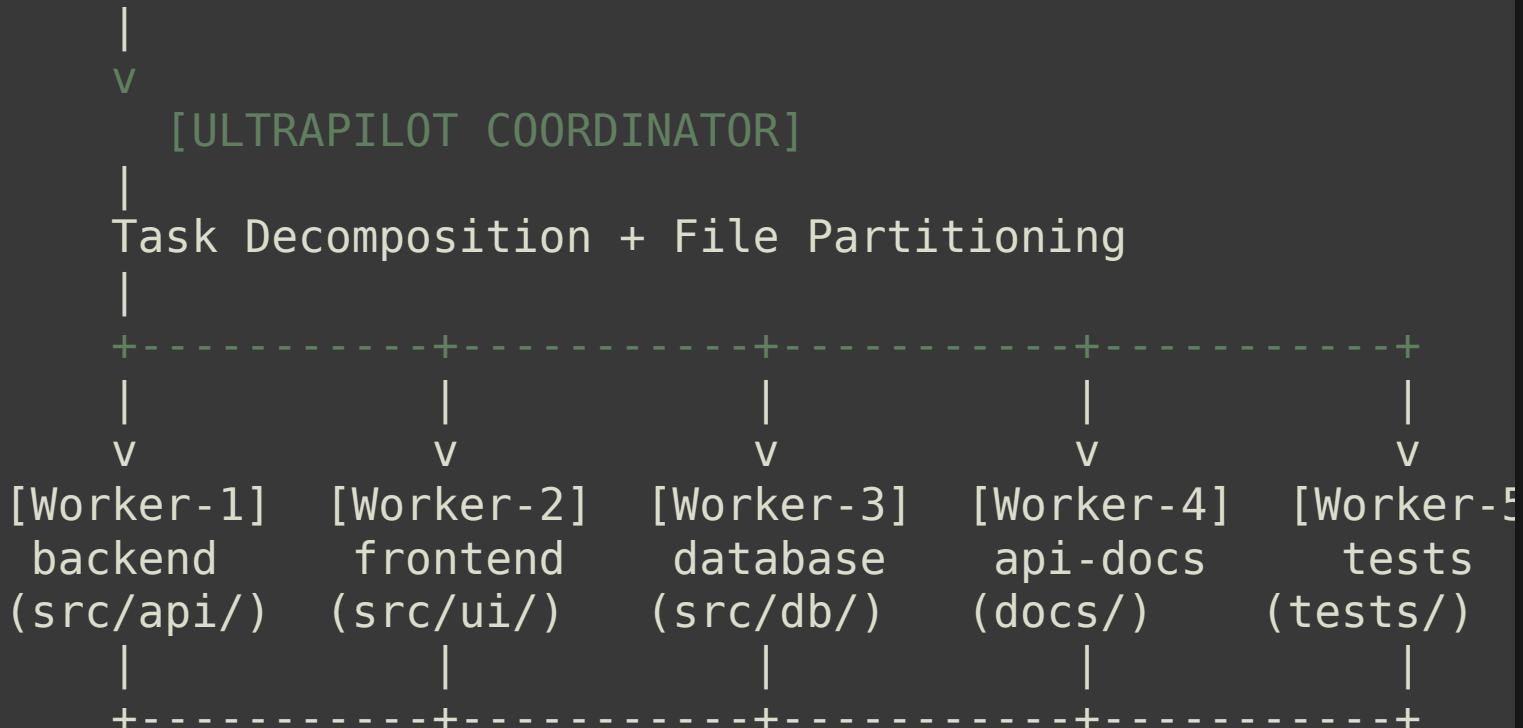
3-5x faster than standard autopilot for suitable tasks.

```
"ultrapilot: build a full-stack todo app"
```

Key Innovation: File ownership partitioning

MODE 2: ULTRAPILOT - HOW IT WORKS

User Input: "Build a full-stack todo app"



MODE 2: ULTRAPILOT - WHEN TO USE IT

Best For:

- Multi-component systems (frontend + backend + database)
- Large refactorings with clear module boundaries
- Multi-service architectures
- Parallel test generation

Speed Comparison:

Task	Autopilot	Ultrapilot
Full-stack app	~75 min	~15 min
Multi-service refactor	~32 min	~8 min
Test coverage	~50 min	~10 min

Trigger:

ultrapilot, parallel build, swarm build

MODE 3: SWARM - WHAT IS IT?

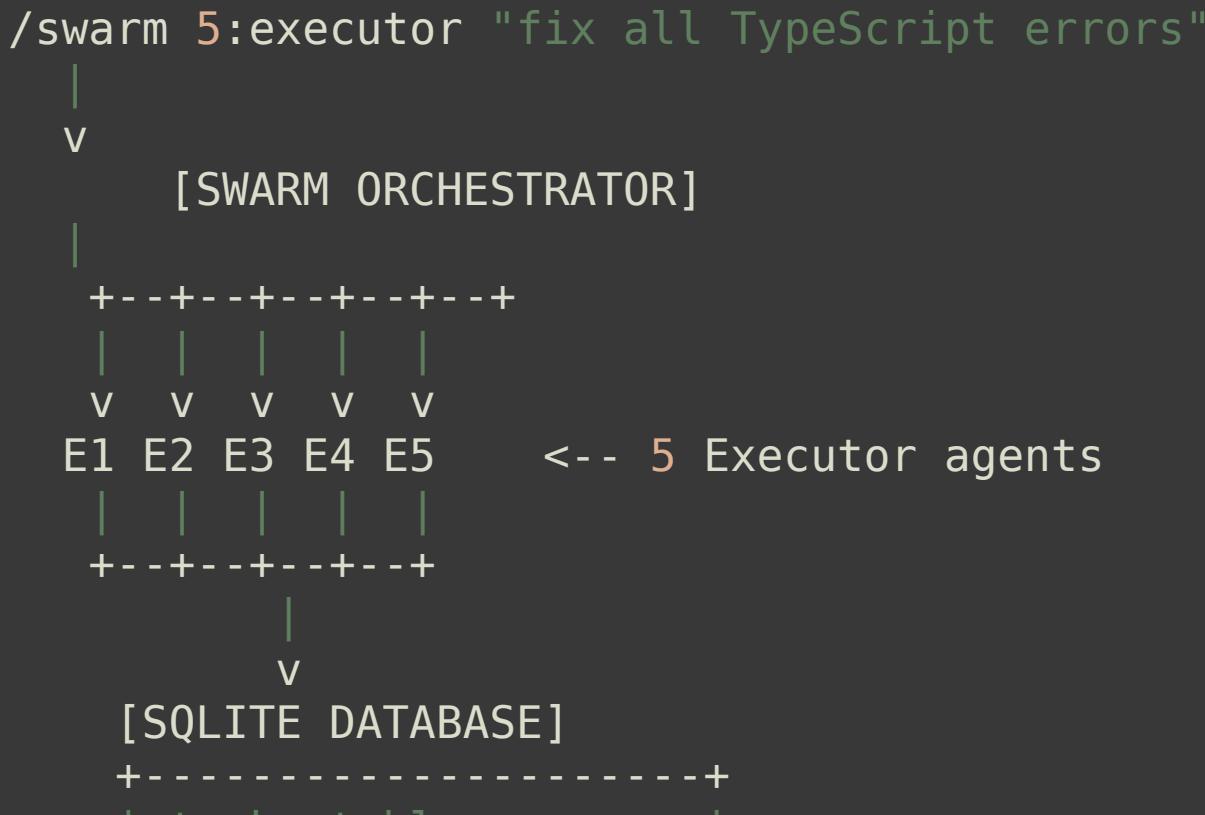
N coordinated agents with atomic task claiming

```
/swarm 5:executor "fix all TypeScript errors"
```

Architecture:

- SQLite-based task pool
- Atomic claiming via transactions
- 5-minute lease timeout with auto-release
- Heartbeat monitoring for fault tolerance

MODE 3: SWARM - HOW IT WORKS



Claim Protocol:

1. Agent calls `claimTask()`
2. SQLite transaction atomically updates status
3. Agent works on task
4. Agent calls `completeTask()` or `failTask()`

MODE 3: SWARM - WHEN TO USE IT

Best For:

- Many independent parallel tasks
- File-by-file operations
- Batch processing

Use Cases:

```
# Fix all TypeScript errors  
/swarm 5:executor "fix all TypeScript errors"  
  
# Style all UI components  
/swarm 3:designer "implement Material-UI styling for all comp
```

```
# Security audit all endpoints  
/swarm 4:security-reviewer "review all API endpoints"  
  
# Add documentation  
/swarm 2:writer "add JSDoc comments to all exported functions"
```

MODE 4: PIPELINE - WHAT IS IT?

Sequential agent chaining with data passing

Like Unix pipes, but for AI agents.

```
/pipeline explore -> architect -> executor "add authentication"
```

Output of one agent becomes input to the next:

```
[explore findings] -> [architect analysis] -> [executor imple
```

MODE 4: PIPELINE - BUILT-IN PRESETS

Preset	Stages	Use For
review	explore -> architect -> critic -> executor	Major features, refactorings
implement	planner -> executor -> tdd-guide	New features with tests

Preset	Stages	Use For
debug	explore -> architect -> build-fixer	Bugs, build errors
research	parallel(researcher, explore) -> architect -> writer	Technology decisions
refactor	explore -> architect-medium -> executor-high -> qa-tester	Safe refactoring
security	explore -> security-reviewer -> executor -> qa-tester	Security fixes

> security-reviewer-low

Usage:

```
/pipeline review "add rate limiting to API"  
/pipeline debug "login fails with OAuth"  
/pipeline security "audit user authentication"
```

MODE 4: PIPELINE - WHEN TO USE IT

Best For:

- Multi-stage processing workflows
- Code review processes
- Research-to-implementation flows

Custom Pipeline Syntax:

```
# Basic sequential  
/pipeline agent1 -> agent2 -> agent3 "task"  
  
# With model specification  
/pipeline explore:haiku -> architect:opus -> executor:sonnet
```

```
# With parallel stages  
/pipeline [explore, researcher] -> architect -> executor "tas
```

Data Flow:

```
{  
  "pipeline_context": {  
    "original_task": "user's request",  
    "previous_stages": [  
      {"agent": "explore", "findings": "..."}  
    ],  
    "current_stage": "architect"  
  }  
}
```

MODE 5: ECOMODE - WHAT IS IT?

Token-efficient parallel execution

30-50% cheaper than standard execution.

```
eco: implement new feature
```

Strategy:

- Prefer Haiku (cheapest) for all tasks
- Only upgrade to Sonnet when needed
- Avoid Opus unless absolutely essential

MODE 5: ECOMODE - HOW IT WORKS

Routing Rules:

Task Type	Standard Mode	Ecomode
Simple lookup	architect-low	architect-low
Standard impl	executor	executor-low (first attempt)

Task Type	Standard Mode	Ecomode	
Complex analysis	architect	architect-medium	
Planning	planner (Opus)	Avoid if possible	
Agent Routing Table:			
Domain	Preferred (Haiku)	Fallback (Sonnet)	Avoid (Opus)
Analysis	architect-	architect-	architect

Domain	Preferred (Haiku)	Fallback (Sonnet)	Avoid (Opus)
	low	medium	
Execution	executor- low	executor	executor- high
Search	explore	explore- medium	explore- high
Frontend	designer- low	designer	designer- high

MODE 5: ECOMODE - WHEN TO USE IT

Best For:

- Budget-conscious projects
- Iterative development (many small changes)
- Exploratory work
- Personal projects

Cost Savings Example:

Task	Standard Cost	Ecomode Cost	Savings
100 simple fixes	~\$3.00	~\$0.50	83%
Feature impl	~\$1.50	~\$0.75	50%
Full build	~\$10.00	~\$5.00	50%

Trigger:

eco, ecomode, efficient, save-tokens, budget

SECTION 3

THE AGENT SYSTEM

32 SPECIALIZED AGENTS

Domain	Agents
Analysis	architect, architect-medium, architect-low
Execution	executor, executor-high, executor-low
Search	explore, explore-medium, explore-high
Research	researcher, researcher-low

Domain Agents

Frontend designer, designer-high, designer-low

Documentation writer

Visual vision

Planning planner, analyst

Critique critic

Testing qa-tester, qa-tester-high

Security security-reviewer, security-reviewer-low

Domain

Agents

Build

build-fixer, build-fixer-low

TDD

tdd-guide, tdd-guide-low

Code Review

code-reviewer, code-reviewer-low

Data Science

scientist, scientist-high, scientist-low

3-TIER MODEL ROUTING

LOW (Haiku)	MEDIUM (Sonnet)	HIGH (Opus)
\$0.25/\$1.25/1M	\$3/\$15/1M	\$15/\$75/1M
Simple lookups	Standard work	Complex reasoning
Quick searches	Feature impl	Architecture
Basic fixes	Moderate debug	Deep debugging
Documentation	UI components	Security audits

Use by default
↓
Only when truly

Upgrade when
↑
Only when truly

Cost Example:

- 1000 simple questions: Haiku = $0.25vs$ *Opus* = 15
(60x cheaper!)

SMART DELEGATION

OMC automatically picks the right agent:

Task	Agent Selected	Model
"What does this function return?"	architect-low	Haiku
"Find where UserService is defined"	explore	Haiku
"Add validation to login form"	executor-low	Haiku

Task	Agent Selected	Model
"Implement OAuth2 flow"	executor	Sonnet
"Debug race condition in auth"	architect	Opus
"Refactor entire auth module"	executor-high	Opus

Delegation Code:

```
Task(  
    subagent_type="oh-my-claudecode:executor-low",  
    model="haiku",
```

AGENT COMPOSITION

Skills + Agents combine for powerful workflows:

```
"ralph ultrawork: migrate database"
|           |
|           +-> Parallel execution (ultrawork)
+-----> Persistence (ralph)
```

Real Example:

```
ralph ultrawork git-master: refactor authentication
|           |           |
|           |           +-> Git expertise (atomic commits)
|           +-----> Maximum parallelism
+-----> Won't stop until verified complete
```

Result: Persistent, parallel, git-aware refactoring

DELEGATION CATEGORIES

Semantic task categorization with auto-detection:

Category	Tier	Temp	Thinking	Auto-Detect
From				
visual-engineering	HIGH	0.7	high	"UI", "comp", "style"
ultrabrain	HIGH	0.3	max	"debug", "archi"

Category	Tier	Temp	Thinking	Auto-Detect
artistry	MEDIUM	0.9	medium	"creative", "talent"
quick	LOW	0.1	low	"find", "discover", "is", "what"
writing	MEDIUM	0.5	medium	"document", "explain"

How It Works:

User: "debug the race condition in auth"

|

v

Detected: "debug" keyword

|

v

Category: ultrabrain

|

v

Settings: HIGH tier, temp=0.3, max thinking

SECTION 4

LIVE DEMO SCENARIOS

DEMO 1: AUTOPILOT

Command:

```
autopilot: build a REST API for a bookstore with CRUD operations
```

What Happens:

1. Expansion Phase (~2 min)

- Analyst extracts: entities (Book, Author), operations (CRUD), constraints
- Architect creates: technical spec, database schema, API design

2. Planning Phase (~1 min)

- Architect creates implementation plan
- Critic validates completeness

3. Execution Phase (~10-15 min)

- Executors implement routes, models, tests in parallel

DEMO 2: ULTRAWORK

Command:

```
ulw fix all TypeScript errors
```

What Happens:

[ULTRAWORK ACTIVATED]

Scanning for TypeScript errors...
Found 23 errors across 8 files.

Spawning parallel agents:

```
[executor-low:1] -> src/api/routes.ts (5 errors)
[executor-low:2] -> src/api/handlers.ts (3 errors)
[executor-low:3] -> src/ui/App.tsx (4 errors)
[executor-low:4] -> src/db/models.ts (6 errors)
[executor-low:5] -> src/utils/helpers.ts (5 errors)
```

Progress: [=====] 100%

All 23 errors fixed in 2m 34s

DEMO 3: PIPELINE

Command:

```
/pipeline review "add rate limiting to the API"
```

What Happens:

PIPELINE: review

Stages: explore -> architect -> critic -> executor

[Stage 1/4: explore] RUNNING

Finding rate limiting patterns in codebase...

Found: middleware pattern in src/middleware/

Found: express-rate-limit in dependencies

OUTPUT: 3 relevant files identified

[Stage 2/4: architect] RUNNING

Analyzing: existing middleware pattern

Recommendation: Use express-rate-limit with Redis store

Design: Per-route configuration

OUTPUT: Implementation plan created

DEMO 4: PLANNING

Command:

```
plan the authentication system
```

What Happens:

[PLAN SKILL ACTIVATED]

I'll help you plan the authentication system. Let me ask a few questions:

Q1: What authentication methods do you need?

- Username/Password
- OAuth2 (Google, GitHub)
- Magic Link (email)
- API Keys

Q2: What session management approach?

- JWT tokens

- [] Server-side sessions
- [] Hybrid

DEMO 5: RALPH

Command:

```
ralph: refactor the auth module to use dependency injection
```

What Happens:

```
[RALPH ACTIVATED - Will not stop until verified complete]
```

Iteration 1/10:

```
Analyzing auth module structure...
```

```
Creating refactoring plan...
```

```
Executing changes...
```

```
ERROR: Test failure in auth.test.ts
```

Iteration 2/10:

```
Analyzing failure: Mock not updated for new DI pattern
```

```
Fixing test mocks...
```

```
Re-running tests...
```

ERROR: Type error in UserService

Iteration 3/10:

SECTION 5

DEVELOPER EXPERIENCE

MAGIC KEYWORDS

Optional shortcuts for power users:

Keyword	Effect	Example
autopilot	Full autonomous execution	autopilot: build todo app
ralph	Persistence until complete	ralph: fix auth bugs
ulw	Maximum parallelism	ulw fix all errors

Keyword	Effect	Example
eco	Token-efficient execution	eco: add validation
plan	Interactive planning	plan the API
ralplan	Iterative planning consensus	ralplan new feature

Combinations work:

```
ralph ulw: migrate database  
^ ^
```

HUD STATUSLINE

Real-time visibility into OMC state:

```
+-----+  
| OMC | autopilot:exec | 3 agents | 5/12 tasks | ctx:45% | $2  
+-----+  
          ^           ^           ^           ^  
          |           |           |           |  
          |           |           |           |  
Active mode    # running    Progress      Context      Cost  
       agents        window
```

Setup:

```
/oh-my-claudecode:hud setup
```

Presets:

- **minimal** - Just active mode
- **focused** - Mode + progress (default)
- **full** - Everything including cost

NOTE PAD WISDOM SYSTEM

Plan-scoped knowledge capture:

Location: .omc/notepads/{plan-name}/

File	Purpose	Example
learnings.md	Technical discoveries	"Redis requires explicit TTL for rate limit keys"
decisions.md	Design decisions	"Chose JWT over sessions for

File

Purpose

Example

issues.md

Known issues

stateless scaling"

problems.md

Blockers

"OAuth callback URL must be HTTPS in prod"

API:

ANALYTICS & COST TRACKING

Track token usage and costs:

```
$ omc-analytics summary

Session Summary (last 7 days)
-----
Total sessions: 23
Total tokens: 1,234,567
Total cost: $18.45

By Model:
Haiku: 890,000 tokens ($0.89)
Sonnet: 300,000 tokens ($4.50)
Opus: 44,567 tokens ($13.06)

By Mode:
autopilot: 45% of cost
            1,088,000 tokens
            $8.27
```

SECTION 6

GETTING STARTED

INSTALLATION

Method 1: Plugin Marketplace (Recommended)

```
/plugin marketplace add https://github.com/Yeachan-Heo/oh-my-  
/plugin install oh-my-claudecode
```

Method 2: NPM Global

```
npm install -g oh-my-claude-sisyphus
```

Method 3: Manual Git Clone

```
git clone https://github.com/Yeachan-Heo/oh-my-claudecode.git  
cd oh-my-claudecode  
npm install && npm run build
```

Requirements:

- Claude Code CLI
- Claude Max/Pro subscription OR Anthropic API key
- Node.js 20+

FIRST STEPS

Step 1: Install

```
/plugin marketplace add https://github.com/Yeachan-Heo/oh-my-  
/plugin install oh-my-claudecode
```

Step 2: Setup

```
/oh-my-claudecode:omc-setup
```

(Configures defaults, HUD, preferences)

Step 3: Build something

```
autopilot: build a REST API for managing tasks
```

That's it. Everything else is automatic.

CONFIGURATION

Project-level: CLAUDE.md in project root Global:
~/.claude/CLAUDE.md

Key Settings:

```
// ~/.claude/settings.json
{
  "omc": {
    "defaultExecutionMode": "ultrawork", // or "ecomode"
    "autopilot": {
      "maxIterations": 10,
      "maxQaCycles": 5,
      "skipValidation": false
    },
    "hud": {
      "preset": "focused"
    }
}
```

```
    }
```

Agent Customization:

- Modify agent prompts in `agents/*.md`
- Override tools per agent
- Create custom agents

SECTION 7

CLOSING

REAL-WORLD USE CASES

Use Case	Best Mode	Why
Backend API development	autopilot	Full end-to-end workflow
Frontend component library	ultrapilot	Many independent components
Database migrations	ralph	Needs persistence

Use Case

Best Mode

Why

CI/CD pipeline setup

pipeline:implement

Sequential stages

Documentation generation

swarm:writer

Parallel doc writing

Bug triage & fixing

swarm:executor

Many independent fixes

Use Case

Best Mode

Why

Security audit

pipeline:security

Structured
review
process

Exploratory
prototyping

ecomode

Budget-
conscious
iteration

RESOURCES

GitHub Repository

github.com/Yeachan-Heo/oh-my-claudecode

Website & Documentation

yeachan-heo.github.io/oh-my-claudecode-website

NPM Package

`npm install -g oh-my-claude-sisyphus`

Documentation Directory

- /docs/REFERENCE.md - Complete feature reference
- /docs/MIGRATION.md - Upgrade guide
- /docs/ARCHITECTURE.md - How it works

Getting Help

- /oh-my-claudecode:help - Usage guide
- /oh-my-claudecode:doctor - Diagnose issues

Q&A

Common Questions:

Question

Does OMC work
with Claude API
keys?

Answer

Yes, both Max/Pro subscription
and API keys work

Can I use OMC with other AI models?

No, OMC is specifically for
Claude Code

Question

Answer

How do I stop a runaway autopilot?

Say "stop", "cancel", or /oh-my-claudecode:cancel

Why is my HUD not showing?

Run /oh-my-claudecode:hud setup

Can I create custom agents?

Yes, add .md files to agents/ directory

Is there a cost limit?

No built-in limit, but ecomode helps control costs

Questions?

THANK YOU

oh-my-claudecode

Zero learning curve. Maximum power.

github.com/Yeachan-Heo/oh-my-claudecode

Get Started Now:

```
/plugin marketplace add https://github.com/Yeachan-Heo/oh-my-  
/plugin install oh-my-claudecode  
autopilot: build something amazing
```

APPENDIX A: COMPLETE AGENT REFERENCE

Agent	Model	Best For
architect	opus	Complex architecture, deep debugging
architect- medium	sonnet	Moderate analysis
architect-low	haiku	Quick code questions
executor	sonnet	Standard implementation

Agent	Model	Best For
executor-high	opus	Complex refactoring
executor-low	haiku	Simple fixes
explore	haiku	Fast file search
explore-medium	sonnet	Pattern matching
explore-high	opus	Architectural search
designer	sonnet	UI components

Agent	Model	Best For
designer-high	opus	Design systems
designer-low	haiku	Simple styling
	--	

APPENDIX A: COMPLETE AGENT REFERENCE (CONTINUED)

Agent	Model	Best For
researcher	sonnet	External docs, APIs

Agent	Model	Best For
researcher-low	haiku	Quick lookups
writer	haiku	Documentation
vision	sonnet	Image analysis
planner	opus	Strategic planning
analyst	opus	Requirements extraction
critic	opus	Plan review
qa-tester	sonnet	CLI testing
qa-tester-high	opus	Comprehensive QA

Agent	Model	Best For
security-reviewer	opus	Security audits
security-reviewer-low	haiku	Quick security scan
	--	

APPENDIX A: COMPLETE AGENT REFERENCE (CONTINUED)

Agent	Model	Best For
build-fixer	sonnet	Build error resolution

APPENDIX B: COMPLETE SKILL REFERENCE

Skill	Purpose	Trigger
autopilot	Full autonomous execution	"autopilot", "build me"
ultrapilot	Parallel autopilot	"ultrapilot", "parallel build"
ralph	Persistence mode	"ralph", "don't stop"

Skill	Purpose	Trigger
ulrawork	Maximum parallelism	"ulw", "ulrawork"
ecomode	Token-efficient mode	"eco", "budget"
swarm	Coordinated agents	/swarm N:agent
pipeline	Sequential chaining	/pipeline preset
plan	Planning interview	"plan the"

Skill	Purpose	Trigger
ralplan	Iterative planning	"ralplan"
cancel	Stop any mode	"stop", "cancel"
	--	

APPENDIX B: COMPLETE SKILL REFERENCE (CONTINUED)

Skill	Purpose	Trigger
analyze	Deep investigation	"analyze", "debug"

Skill	Purpose	Trigger
deepsearch	Thorough search	"search", "find"
deepinit	Generate AGENTS.md	"index codebase"
frontend-ui-ux	Design sensibility	UI context (auto)
git-master	Git expertise	Git context (auto)
ultraqa	QA cycling	"test", "QA"
learner	Extract skills	"extract skill"

Skill	Purpose	Trigger
note	Save to notepad	"remember", "note"
hud	Configure HUD	/hud
doctor	Diagnose issues	/doctor
	--	

APPENDIX B: COMPLETE SKILL REFERENCE (CONTINUED)

Skill	Purpose	Trigger
help	Show usage guide	/help
omc-setup	Setup wizard	/omc-setup
ralph-init	Initialize PRD	/ralph-init
release	Release workflow	/release
review	Review plan	"review plan"
research	Scientist orchestration	"research", "statistics"
tdd	TDD enforcement	"tdd", "test first"

APPENDIX C: KEYBOARD SHORTCUTS SUMMARY

Shortcut	Full Command	Effect
autopilot:	/oh-my-claudecode:autopilot	Full autonode mode
ralph:	/oh-my-claudecode:ralph	Persistent mode
ulw	/oh-my-claudecode:ultrawork	Parallel execution

Shortcut

Full Command

Effect

eco:

/oh-my-claudecode:ecomode

Token-efficient mode

plan

/oh-my-claudecode:plan

Planning interview

Combinations:

```
ralph ulw: task          # Persistent + Parallel  
ralph eco: task          # Persistent + Efficient  
autopilot eco: task      # Auto + Efficient (eco wins)
```